

## Steps to Identify and Mitigate High CPU SQL Server Usage Queries

### Problem Summary

- SQL Server can consume high CPU, leading to poor performance.
- Symptoms include:
  - Slow simple queries
  - Lagging reports
  - CPU process (sqlservr.exe) near 100%
- High CPU is commonly related to *query patterns, missing indexes, or out-of-date statistics*.

### Causes of High CPU Usage

Common contributors:

- High logical reads from missing indexes or stale statistics
- Sudden workload spike
- Concurrency
- Parallel query execution
- Long-running queries
- Heavy calculations
- Sorts, Temp Tables
- Other applications competing for CPU
- Compression tasks, backups, etc.

### Simulating CPU Load

To generate CPU load for demonstration (using SQLQueryStress):

```
SELECT TOP (1000)
[WorkOrderID], [ProductID], [OrderQty], [StockedQty],
[ScrappedQty], [StartDate], [EndDate], [DueDate],
[ScrapReasonID], [ModifiedDate]
FROM [AdventureWorks2022].[Production].[WorkOrder]
WHERE OrderQty = 5;
GO
```

```
SELECT TOP (1000)
[BusinessEntityID], [PersonType], [NameStyle], [Title],
[FirstName], [MiddleName], [LastName], [Suffix],
[EmailPromotion], [AdditionalContactInfo], [Demographics],
[rowguid], [ModifiedDate]
FROM [AdventureWorks2022].[Person].[Person]
WHERE FirstName = 'ken';
GO
```

*Used to stress the server CPU.*

### Checking Overall CPU Usage

- Task Manager can show overall and SQL Server CPU utilization.
- If SQL Server engine process is using large percentage of CPU, proceed with deeper analysis.

## Identify Top CPU-Consuming Queries

Use DMVs to find queries consuming the most CPU:

```

SELECT TOP 10
    COALESCE(DB_NAME(st.dbid),DB_NAME(CONVERT(INT, qp.dbid))) AS [DatabaseName],
    qs.creation_time AS [PlanCreationTime],
    qs.last_execution_time AS [LastExecutedTime],
    qs.execution_count AS [ExecutionCount],
    qs.total_worker_time / 1000 AS [Total_CPU_Time_ms],
    (qs.total_worker_time / qs.execution_count) / 1000 AS [Avg_CPU_Time_ms],
    (qs.total_worker_time / 1000) AS [Cumulative_CPU_Time_All_Executions_ms],
    (qs.total_logical_reads + qs.total_logical_writes) AS [TotalLogicalIO],
    (qs.total_logical_reads + qs.total_logical_writes) / qs.execution_count AS [AvgLogicalIO],
    SUBSTRING(st.text,
        (qs.statement_start_offset / 2) + 1,
        ((CASE qs.statement_end_offset WHEN -1 THEN DATALENGTH(st.text)
        ELSE qs.statement_end_offset END - qs.statement_start_offset)/2)+1) AS [QueryText],
    qp.query_plan AS [ExecutionPlan]
FROM sys.dm_exec_query_stats AS qs
CROSS APPLY sys.dm_exec_sql_text(qs.sql_handle) AS st
CROSS APPLY sys.dm_exec_query_plan(qs.plan_handle) AS qp
WHERE COALESCE(DB_NAME(st.dbid),DB_NAME(CONVERT(INT, qp.dbid)))
    NOT IN ('master','model','msdb','tempdb','DBADashDB')
ORDER BY qs.total_worker_time DESC;
Finds top CPU-heavy queries.
```

## Queries with Missing Indexes

Missing indexes can cause scans instead of seeks, increasing CPU:

```

WITH XMLNAMESPACES ('http://schemas.microsoft.com/sqlserver/2004/07/showplan' AS p)
SELECT
    COALESCE(DB_NAME(st.dbid), DB_NAME(CONVERT(INT, qp.dbid))) AS DatabaseName,
    qs.creation_time AS PlanCreationTime,
    qs.last_execution_time AS LastExecutedTime,
    qs.execution_count AS ExecutionCount,
    qs.total_worker_time / 1000 AS Total_CPU_Time_ms,
    (qs.total_worker_time / qs.execution_count) / 1000 AS Avg_CPU_Time_ms,
    (qs.total_worker_time / 1000) AS Cumulative_CPU_Time_All_Executions_ms,
    (qs.total_logical_reads + qs.total_logical_writes) AS TotalLogicalIO,
    (qs.total_logical_reads + qs.total_logical_writes) / qs.execution_count AS AvgLogicalIO,
    SUBSTRING(st.text,
        (qs.statement_start_offset/2)+1,
        ((CASE qs.statement_end_offset WHEN -1 THEN DATALENGTH(st.text)
        ELSE qs.statement_end_offset END - qs.statement_start_offset)/2)+1) AS [QueryText],
    qp.query_plan AS [ExecutionPlan]
```

```

FROM (
    SELECT TOP 10 *
    FROM sys.dm_exec_query_stats
    ORDER BY total_worker_time DESC
) AS qs
CROSS APPLY sys.dm_exec_sql_text(qs.sql_handle) AS st
CROSS APPLY sys.dm_exec_query_plan(qs.plan_handle) AS qp
WHERE qp.query_plan.value('count(/p:MissingIndexGroup)', 'int') > 0;
GO
Lists top CPU queries with missing index hints in execution plans.

```

### Script to Generate Missing Index Create Statements

```

SELECT TOP 25
    DB_NAME(dm_mid.database_id) AS DatabaseID,
    dm_migs.avg_user_impact*(dm_migs.user_seeks+dm_migs.user_scans) Avg_Estimated_Impact,
    dm_migs.last_user_seek AS Last_User_Seek,
    OBJECT_NAME(dm_mid.OBJECT_ID,dm_mid.database_id) AS [TableName],
    'CREATE INDEX [IX_'
    + OBJECT_NAME(dm_mid.OBJECT_ID,dm_mid.database_id)
    + '_' + REPLACE(REPLACE(REPLACE(ISNULL(dm_mid.equality_columns,'','_'),',','_'),'[',''),']','')
    + CASE WHEN dm_mid.equality_columns IS NOT NULL
        AND dm_mid.inequality_columns IS NOT NULL THEN ' ' ELSE '' END
    + REPLACE(REPLACE(REPLACE(ISNULL(dm_mid.inequality_columns,'','_'),',','_'),'[',''),']','')
    + '] ON ' + dm_mid.statement
    + ' (' + ISNULL(dm_mid.equality_columns,'')
    + CASE WHEN dm_mid.equality_columns IS NOT NULL
        AND dm_mid.inequality_columns IS NOT NULL THEN ' ' ELSE '' END
    + ISNULL(dm_mid.inequality_columns,'')
    + ')'
    + ISNULL(' INCLUDE (' + dm_mid.included_columns + ')', '') AS Create_Statement
FROM sys.dm_db_missing_index_groups dm_mig
JOIN sys.dm_db_missing_index_group_stats dm_migs
    ON dm_migs.group_handle = dm_mig.index_group_handle
JOIN sys.dm_db_missing_index_details dm_mid
    ON dm_mig.index_handle = dm_mid.index_handle
WHERE dm_mid.database_ID = DB_ID()
ORDER BY Avg_Estimated_Impact DESC;
GO
Builds recommended index scripts based on missing index DMVs.

```

**Important:** Before creating indexes, evaluate impact — unnecessary or redundant indexes can degrade performance.

## Finding Out-of-Date Statistics

Outdated statistics can lead to poor query plans:

```

SELECT DISTINCT
    OBJECT_NAME(s.[object_id]) AS TableName,
    c.name AS ColumnName,
    s.name AS StatName,
    STATS_DATE(s.[object_id], s.stats_id) AS LastUpdated,
    DATEDIFF(d, STATS_DATE(s.[object_id], s.stats_id), GETDATE()) AS DaysOld,
    dsp.modification_counter,
    s.auto_created,
    s.user_created,
    s.no_recompute,
    s.[object_id],
    s.stats_id,
    sc.stats_column_id,
    sc.column_id
FROM sys.stats s
JOIN sys.stats_columns sc
    ON sc.[object_id] = s.[object_id]
    AND sc.stats_id = s.stats_id
JOIN sys.columns c
    ON c.[object_id] = sc.[object_id]
    AND c.column_id = sc.column_id
JOIN sys.partitions par
    ON par.[object_id] = s.[object_id]
JOIN sys.objects obj
    ON par.[object_id] = obj.[object_id]
CROSS APPLY sys.dm_db_stats_properties(sc.[object_id], s.stats_id) AS dsp
WHERE OBJECTPROPERTY(s.OBJECT_ID,'IsUserTable') = 1
    AND (s.auto_created = 1 OR s.user_created = 1)
ORDER BY DaysOld;
GO

```

*Shows age of statistics per object/column.*

## Update Statistics

If outdated, update:

EXEC sp\_updatestats;

*Refreshes statistics across user tables.*

⚠ This can force recompiles and temporarily increase load.

**Step-by-step operational playbook** we can actually use in production to **diagnose, analyze, and fix High CPU issues in SQL Server**, based on real-world DBA practice.

## SQL Server High CPU – Step-by-Step Troubleshooting Playbook

### STEP 1: Confirm SQL Server Is the CPU Culprit

#### OS Level

- Open **Task Manager**
- Check:
  - Overall CPU %
  - sqlservr.exe CPU usage
- If SQL Server is consuming most CPU → proceed

#### SQL Level (optional quick check)

SELECT

```
cpu_count,
scheduler_count,
hyperthread_ratio
FROM sys.dm_os_sys_info;
```

### STEP 2: Identify Top CPU-Consuming Queries (MOST IMPORTANT STEP)

Run this during high CPU

```
SELECT TOP 10
COALESCE(DB_NAME(st.dbid),DB_NAME(CONVERT(INT, qp.dbid))) AS DatabaseName,
qs.execution_count,
qs.total_worker_time / 1000 AS Total_CPU_ms,
(qs.total_worker_time / qs.execution_count) / 1000 AS Avg_CPU_ms,
qs.total_logical_reads,
qs.last_execution_time,
SUBSTRING(st.text,
(qs.statement_start_offset/2)+1,
((CASE qs.statement_end_offset
WHEN -1 THEN DATALENGTH(st.text)
ELSE qs.statement_end_offset END
- qs.statement_start_offset)/2)+1) AS QueryText,
qp.query_plan
FROM sys.dm_exec_query_stats qs
CROSS APPLY sys.dm_exec_sql_text(qs.sql_handle) st
CROSS APPLY sys.dm_exec_query_plan(qs.plan_handle) qp
ORDER BY qs.total_worker_time DESC;
```

#### How to interpret

- **High Total CPU** → long-running / heavy queries
- **High Avg CPU** → inefficient logic
- **High execution\_count + medium CPU** → “death by 1000 cuts”

**Action:** Copy top 3 queries and focus only on those

### STEP 3: Check for Missing Indexes in CPU-Heavy Queries

```
WITH XMLNAMESPACES ('http://schemas.microsoft.com/sqlserver/2004/07/showplan' AS p)
SELECT
    qs.total_worker_time / 1000 AS Total_CPU_ms,
    qs.execution_count,
    SUBSTRING(st.text,
        (qs.statement_start_offset/2)+1,
        ((CASE qs.statement_end_offset
            WHEN -1 THEN DATALENGTH(st.text)
            ELSE qs.statement_end_offset END
        - qs.statement_start_offset)/2)+1) AS QueryText,
    qp.query_plan
FROM sys.dm_exec_query_stats qs
CROSS APPLY sys.dm_exec_sql_text(qs.sql_handle) st
CROSS APPLY sys.dm_exec_query_plan(qs.plan_handle) qp
WHERE qp.query_plan.value('count(/p:MissingIndexGroup)', 'int') > 0
ORDER BY qs.total_worker_time DESC;
```

#### What to look for

- MissingIndexGroup in execution plan
- Table scans instead of index seeks

### STEP 4: Generate Missing Index Scripts (DO NOT AUTO-CREATE)

```
SELECT TOP 25
    DB_NAME(dm_mid.database_id) AS DatabaseName,
    dm_migs.avg_user_impact*(dm_migs.user_seeks+dm_migs.user_scans) AS Impact,
    OBJECT_NAME(dm_mid.OBJECT_ID,dm_mid.database_id) AS TableName,
    'CREATE INDEX IX_'
    + OBJECT_NAME(dm_mid.OBJECT_ID,dm_mid.database_id)
    + ' ON ' + dm_mid.statement
    + ' (' + ISNULL(dm_mid.equality_columns,'')
    + CASE WHEN dm_mid.equality_columns IS NOT NULL
        AND dm_mid.inequality_columns IS NOT NULL THEN ',' ELSE " "
        END
    + ISNULL(dm_mid.inequality_columns,'')
    + ')'
    + ISNULL(' INCLUDE (' + dm_mid.included_columns + ')', '') AS CreateIndexScript
FROM sys.dm_db_missing_index_groups dm_mig
JOIN sys.dm_db_missing_index_group_stats dm_migs
    ON dm_migs.group_handle = dm_mig.index_group_handle
JOIN sys.dm_db_missing_index_details dm_mid
    ON dm_mig.index_handle = dm_mid.index_handle
WHERE dm_mid.database_id = DB_ID()
ORDER BY Impact DESC;
```

#### Before creating index

Check:

- Existing similar indexes
- Write workload (INSERT/UPDATE heavy?)
- Index size & maintenance impact

Never blindly create all suggested indexes

## STEP 5: Check Statistics Age & Modifications

```

SELECT
    OBJECT_NAME(s.object_id) AS TableName,
    s.name AS StatName,
    STATS_DATE(s.object_id, s.stats_id) AS LastUpdated,
    dsp.modification_counter,
    DATEDIFF(day, STATS_DATE(s.object_id, s.stats_id), GETDATE()) AS DaysOld
FROM sys.stats s
CROSS APPLY sys.dm_db_stats_properties(s.object_id, s.stats_id) dsp
WHERE OBJECTPROPERTY(s.object_id, 'IsUserTable') = 1
ORDER BY DaysOld DESC;

```

### Red flags

- Stats > 7–30 days old
- High modification\_counter

## STEP 6: Update Statistics (Controlled)

### Safer (table-specific)

```
UPDATE STATISTICS dbo.TableName WITH FULLSCAN;
```

### Broad (use carefully)

```
EXEC sp_updatestats;
```

Expect recompiles & short CPU spike

## STEP 7: Look for Parallelism Issues

### Check CXPACKET / CXCONSUMER waits

```

SELECT *
FROM sys.dm_os_wait_stats
WHERE wait_type LIKE 'CX%';

```

### Check MAXDOP

```

SELECT value_in_use
FROM sys.configurations
WHERE name = 'max degree of parallelism';

```

If misconfigured → CPU spikes due to excessive parallelism

## STEP 8: Query Design Red Flags (Very Common)

Check if top queries have:

- SELECT \*
- Functions in WHERE clause
- Implicit conversions
- Scalar UDFs
- Large sorts (ORDER BY)
- Cartesian joins
- Parameter sniffing symptoms

## STEP 9: Use Query Store (If Enabled)

```

SELECT *
FROM sys.database_query_store_options;

```

If ON:

- Identify regressed plans
- Force good plan
- Compare CPU usage historically

#### STEP 10: Validate Fix

After changes:

1. Re-run **Top CPU query DMV**
2. Compare:
  - o Total CPU
  - o Avg CPU
  - o Logical reads
3. Monitor CPU trend for 24 hours

#### DBA Rule of Thumb

Symptom	Likely Fix
High CPU + scans	Missing index
High CPU + low reads	Bad logic
Sudden CPU spike	Stats or plan change
CPU only during peak	Concurrency / parallelism
CPU steady always	Inefficient design

<https://www.sqlbachamps.com/>

## SQL Server High CPU – COMPLETE END-TO-END GUIDE

(Troubleshooting • Execution Plans • Query Tuning • Emergency Actions • Production Checklist)

### PART 1: IMMEDIATE “CPU 100%” EMERGENCY ACTIONS

When users are screaming and CPU is pegged

#### 1) Confirm SQL Server is the cause

- Task Manager → sqlservr.exe high CPU
- If not SQL Server → stop here

#### 2) Identify what is running RIGHT NOW

```
SELECT
    r.session_id,
    r.status,
    r.cpu_time,
    r.total_elapsed_time,
    r.logical_reads,
    SUBSTRING(t.text, r.statement_start_offset/2,
        (CASE WHEN r.statement_end_offset = -1
            THEN LEN(t.text)
            ELSE r.statement_end_offset END - r.statement_start_offset)/2) AS RunningQuery
FROM sys.dm_exec_requests r
CROSS APPLY sys.dm_exec_sql_text(r.sql_handle) t
ORDER BY r.cpu_time DESC;
```

If one query dominates CPU, that's your target.

#### 3) Emergency kill (LAST RESORT)

KILL <session\_id>;

Use only if business impact is severe.

### PART 2: FIND TOP CPU QUERIES (HISTORICAL)

```
SELECT TOP 10
    DB_NAME(st.dbid) AS DatabaseName,
    qs.execution_count,
    qs.total_worker_time / 1000 AS Total_CPU_ms,
    (qs.total_worker_time / qs.execution_count) / 1000 AS Avg_CPU_ms,
    qs.total_logical_reads,
    qs.last_execution_time,
    SUBSTRING(st.text,
        (qs.statement_start_offset/2)+1,
        ((CASE qs.statement_end_offset
            WHEN -1 THEN DATALENGTH(st.text)
            ELSE qs.statement_end_offset END
            - qs.statement_start_offset)/2)+1) AS QueryText,
    qp.query_plan
FROM sys.dm_exec_query_stats qs
CROSS APPLY sys.dm_exec_sql_text(qs.sql_handle) st
CROSS APPLY sys.dm_exec_query_plan(qs.plan_handle) qp
ORDER BY qs.total_worker_time DESC;
```

### How to read results

Pattern	Meaning
High Total CPU	Heavy query
High Avg CPU	Inefficient logic
High executions	Chatty app
High reads	Missing index

### PART 3: EXECUTION PLAN – HOW TO READ (VERY IMPORTANT)

Focus ONLY on these operators

Operator	Meaning
Table Scan	No useful index
Index Scan	Partial index issue
Index Seek	GOOD
Hash Match	CPU-heavy joins
Sort	CPU + memory heavy
Key Lookup	Missing INCLUDE

### Red flags

- Scan on large table
- Sort + Hash Match together
- Missing Index warnings
- Estimated vs Actual rows mismatch

### PART 4: REAL QUERY TUNING EXAMPLE

#### Bad Query

```
SELECT *
FROM Orders
WHERE YEAR(OrderDate) = 2024;
```

#### Why bad?

- Function on column
- Forces table scan
- High CPU

#### Fixed Query

```
SELECT *
FROM Orders
WHERE OrderDate >= '2024-01-01'
    AND OrderDate < '2025-01-01';
```

- Allows index seek
- CPU drops massively

### PART 5: MISSING INDEX DETECTION

WITH XMLNAMESPACES ('http://schemas.microsoft.com/sqlserver/2004/07/showplan' AS p)

SELECT

```

qs.total_worker_time / 1000 AS Total_CPU_ms,
SUBSTRING(st.text,
  (qs.statement_start_offset/2)+1,
  ((CASE qs.statement_end_offset
    WHEN -1 THEN DATALENGTH(st.text)
    ELSE qs.statement_end_offset END
    - qs.statement_start_offset)/2)+1) AS QueryText,
qp.query_plan
FROM sys.dm_exec_query_stats qs
CROSS APPLY sys.dm_exec_sql_text(qs.sql_handle) st
CROSS APPLY sys.dm_exec_query_plan(qs.plan_handle) qp
WHERE qp.query_plan.value('count(/p:MissingIndexGroup)', 'int') > 0
ORDER BY qs.total_worker_time DESC;

```

## PART 6: GENERATE INDEX SCRIPTS (REVIEW FIRST)

```

SELECT TOP 25
  DB_NAME(dm_mid.database_id) AS DatabaseName,
  dm_migs.avg_user_impact*(dm_migs.user_seeks+dm_migs.user_scans) AS Impact,
  OBJECT_NAME(dm_mid.object_id, dm_mid.database_id) AS TableName,
  'CREATE INDEX IX_' + OBJECT_NAME(dm_mid.object_id, dm_mid.database_id)
  + ' ON ' + dm_mid.statement
  + ' (' + ISNULL(dm_mid.equality_columns,'')
  + CASE WHEN dm_mid.equality_columns IS NOT NULL
    AND dm_mid.inequality_columns IS NOT NULL THEN ',' ELSE '' END
  + ISNULL(dm_mid.inequality_columns,'') + ')'
  + ISNULL(' INCLUDE (' + dm_mid.included_columns + ')','') AS CreateIndexScript
FROM sys.dm_db_missing_index_details dm_mid
JOIN sys.dm_db_missing_index_groups dm_mig
  ON dm_mid.index_handle = dm_mig.index_handle
JOIN sys.dm_db_missing_index_group_stats dm_migs
  ON dm_mig.index_group_handle = dm_migs.group_handle
ORDER BY Impact DESC;
Never auto-create all indexes

```

## PART 7: STATISTICS ANALYSIS

```

SELECT
  OBJECT_NAME(s.object_id) AS TableName,
  s.name AS StatName,
  STATS_DATE(s.object_id, s.stats_id) AS LastUpdated,
  dsp.modification_counter,
  DATEDIFF(day, STATS_DATE(s.object_id, s.stats_id), GETDATE()) AS DaysOld
FROM sys.stats s
CROSS APPLY sys.dm_db_stats_properties(s.object_id, s.stats_id) dsp
WHERE OBJECTPROPERTY(s.object_id,'IsUserTable') = 1
ORDER BY DaysOld DESC;
Update stats

```

```
EXEC sp_updatestats;
or
UPDATE STATISTICS dbo.TableName WITH FULLSCAN;
```

#### PART 8: PARALLELISM & MAXDOP

```
SELECT *
FROM sys.dm_os_wait_stats
WHERE wait_type LIKE 'CX%';
SELECT value_in_use
FROM sys.configurations
WHERE name = 'max degree of parallelism';
```

Bad MAXDOP = CPU chaos

#### PART 9: COMMON CPU KILLERS (MEMORIZE)

- X SELECT \*
- X Scalar UDFs
- X Implicit conversions
- X Cursors
- X RBAR logic
- X Large sorts
- X Parameter sniffing

<https://www.sqlbachamps.com/>

#### PART 10: PRODUCTION CPU CHECKLIST (PRINT THIS)

- ✓ Identify top CPU queries
- ✓ Check execution plan
- ✓ Fix scans → seeks
- ✓ Add **correct** indexes
- ✓ Update statistics
- ✓ Validate MAXDOP
- ✓ Retest under load
- ✓ Monitor 24 hours

#### GOLDEN DBA RULE

90% of high CPU issues are fixed by:

- Indexes
- Statistics
- Query rewrites

## SQL SERVER CPU & PERFORMANCE — RUNBOOK

### 1. PRODUCTION QUERY TUNING (NEED QUERY)

**What is needed:**

1. Query text
2. Actual execution plan (XML or screenshot summary)
3. Table sizes (approx rows)
4. Indexes on involved tables
5. Parameter values used during slow run

**What we have to do:**

- Identify **CPU driver operator**
- Check **row estimate vs actual**
- Rewrite predicates (SARGability)
- Remove unnecessary operators
- Design **minimal correct index**
- Validate with logical reads + CPU delta

### 2. EXECUTION PLAN DECISION TREES (DBA BRAIN MAP)

**ROOT QUESTION**

**What operator consumes the most % cost or CPU?**

#### A. TABLE / INDEX SCAN

Scan?

- └ Small table (<10k rows)? → OK
- └ Large table?
  - └ Predicate non-SARGable? → Rewrite WHERE clause
  - └ Missing index? → Create index
  - └ Wrong index chosen?
    - └ Stats outdated → Update stats
    - └ Parameter sniffing → Fix sniffing

#### B. HASH MATCH (JOIN / AGGREGATE)

Hash Match?

- └ Large input? → Reduce rows earlier
- └ Missing join index? → Add index on join column
- └ Memory spill? → Increase memory / rewrite
- └ Should be Nested Loop? → Fix estimates

#### C. SORT

Sort?

- └ ORDER BY? → Index with correct key order
- └ DISTINCT? → Remove if unnecessary
- └ GROUP BY? → Covering index
- └ Spill to tempdb? → Memory pressure

## D. KEY LOOKUP

Key Lookup?

- └─ High executions? → INCLUDE columns
- └─ Few executions? → Ignore

## E. NESTED LOOP (HIGH CPU)

Nested Loop?

- └─ Inner input large? → Missing index
- └─ Loop count high? → Join order problem
- └─ Parameter sniffing? → Fix sniffing

## 3. PARAMETER SNIFFING — DEEP DIVE (VERY IMPORTANT)

**What it is**

SQL Server:

1. Compiles plan using **first parameter values**
2. Reuses that plan for **all future executions**
3. Disaster when data is skewed

### Classic Example

```
CREATE PROC GetOrders @CustomerID INT
```

AS

SELECT \*

FROM Orders

WHERE CustomerID = @CustomerID;

- CustomerID = 1 → 5 rows
- CustomerID = 999 → 5 million rows

⚠ Same plan used for both

### How to DETECT sniffing

- Query fast sometimes, slow other times
- Execution plan flips between Scan/Seek
- Estimated vs Actual rows wildly different
- Clearing cache temporarily “fixes” it

DBCC FREEPROCCACHE; -- NEVER in prod without approval

## FIX OPTIONS (IN ORDER OF SAFETY)

### Option 1: OPTIMIZE FOR UNKNOWN (BEST DEFAULT)

WHERE CustomerID = @CustomerID

OPTION (OPTIMIZE FOR UNKNOWN);

### Option 2: RECOMPILE (CPU COST)

OPTION (RECOMPILE);

✓ Fixes sniffing

✗ Increases compile CPU

### Option 3: Local Variable Trick

```
DECLARE @CID INT = @CustomerID;
```

```
WHERE CustomerID = @CID;
```

### Option 4: Dynamic SQL (ADVANCED)

Use when query shape must change based on parameter

#### What NOT to do

- ✗ Blanket WITH RECOMPILE everywhere
- ✗ Clearing plan cache regularly
- ✗ Ignoring skewed data distributions

## 4. QUERY STORE & FORCED PLANS — DEEP EXPLANATION

#### What Query Store Does

- Captures **query text**
- Stores **multiple execution plans**
- Tracks **runtime stats over time**
- Lets you **force a known-good plan**

#### When to use Forced Plans

Plan regression after:

- Index change
- Stats update
- SQL upgrade
- Parameter sniffing fix failed

Do NOT use to hide bad queries

#### How to Identify Regression

- Same query
- New plan
- Higher CPU / duration

```
SELECT *
```

```
FROM sys.query_store_runtime_stats
ORDER BY avg_cpu_time DESC;
```

#### Force a Plan

```
EXEC sp_query_store_force_plan
@query_id = 123,
@plan_id = 456;
```

#### What Happens When Forced

- SQL Server **must use that plan**
- If plan becomes invalid → force fails
- SQL Server logs force failures

### Forced Plan Risks

- ⚠ Schema change breaks plan
- ⚠ Index dropped → failure
- ⚠ Data distribution shifts → bad performance

Forced plans are **stabilizers**, not fixes.

### Best Practice

1. Tune query first
2. Fix indexes & stats
3. Use forced plan **only if regression persists**

## 5. CPU INCIDENT RUNBOOK (COPY / PASTE)

### INCIDENT START

- Confirm sqlservr.exe CPU
- Capture top CPU sessions
- Save execution plans
- Identify top operator

### ANALYSIS

- Scan vs Seek
- Row estimate accuracy
- Missing index?
- Parameter sniffing?

### REMEDIATION

- Rewrite query
- Add index (minimal)
- Update stats
- Adjust MAXDOP if needed

### STABILIZATION

- Query Store verify
- Force plan if regression
- Monitor 24 hours