

Preventing Deadlocks in SQL Server:

Deadlocks in **SQL Server** can cause serious performance issues if not understood and prevented properly. Let us see how to prevent them.

1. Understand Deadlocks

- A **deadlock** occurs when **two or more sessions block each other**, each holding a lock and waiting for the other to release a resource.
- SQL Server automatically detects deadlocks and **chooses a victim** (terminates one session).
- Common causes:
 - Accessing resources in different orders.
 - Long transactions with multiple locks.
 - Missing indexes → table scans & unnecessary locks.
 - Poor query design.

2. Best Practices to Prevent Deadlocks

✓ A. Query & Transaction Design

1. **Keep transactions short and simple**
 - The longer a transaction runs, the higher the chance of deadlocks.
 - Avoid holding transactions open while waiting for user input.
2. **Access objects in a consistent order**
 - If all applications access tables in the **same sequence**, circular locking is avoided.
 - Example: Always update TableA before TableB.
3. **Use proper isolation levels**
 - Default READ COMMITTED is fine for many workloads.
 - Use **READ COMMITTED SNAPSHOT (RCSI)** or **SNAPSHOT isolation** to reduce shared locks.
 - Be cautious: SNAPSHOT uses versioning in tempdb.
4. **Avoid unnecessary locking hints**
 - Don't force HOLDLOCK or TABLOCKX unless absolutely needed.

✓ B. Indexing Strategy

1. **Create appropriate indexes**
 - Missing indexes → SQL Server does table scans → more locks.
 - Well-designed indexes minimize locking scope.
2. **Covering indexes**
 - Queries that hit covering indexes avoid touching the base table → fewer lock conflicts.
3. **Filtered indexes**
 - Reduce lock contention by narrowing focus to selective rows.

✓ C. Code & Application Layer

1. **Break large transactions**
 - Commit in batches instead of one big transaction (e.g., 10,000 rows at a time).
2. **Use retry logic in applications**
 - If deadlock happens, catch error 1205 and retry automatically.
3. **Avoid user interaction inside transactions**
 - Don't pause a transaction to wait for user confirmation.
4. **Access fewer rows**
 - Use WHERE clauses properly.
 - Avoid SELECT * → reduce lock footprint.

✅ D. Locking & Concurrency Control

1. **Row versioning**
 - Enable **RCSI** (Read Committed Snapshot Isolation) at the database level:
 - `ALTER DATABASE MyDB SET READ_COMMITTED_SNAPSHOT ON;`
 - This allows readers to use row versions instead of blocking writers.
2. **Use NOLOCK carefully**
 - `WITH (NOLOCK)` avoids shared locks but allows dirty reads.
 - Better approach: RCSI instead of NOLOCK.
3. **Apply UPDLOCK + HOLDLOCK cautiously**
 - Can prevent deadlocks by serializing access — but may hurt concurrency.

3. Monitoring & Troubleshooting Deadlocks

1. **Enable Deadlock Trace Flags**
 - `DBCC TRACEON (1222, -1);` logs deadlock graphs to SQL Server error log.
2. **Use Extended Events / Profiler**
 - Capture `xml_deadlock_report` event.
 - Deadlock graphs show victim, resource, and processes involved.
3. **Query DMVs**
 - Recent deadlocks:
 - `SELECT * FROM sys.dm_tran_locks;`
 - `SELECT * FROM sys.dm_exec_requests;`
4. **Analyze blocking chains**
 - Use `sp_whoisactive` or `sys.dm_exec_requests` to see blockers.

4. Example of Preventing Deadlocks

Problem:

Two transactions:

- Transaction 1 updates TableA then TableB.
 - Transaction 2 updates TableB then TableA.
- This creates a cycle → deadlock.

Fix:

- Enforce consistent order: Always update TableA first, then TableB.

5. Summary of Deadlock Prevention

- ✅ Keep transactions short.
- ✅ Access resources in consistent order.
- ✅ Use proper indexing.
- ✅ Enable RCSI for readers.
- ✅ Avoid unnecessary locks.
- ✅ Monitor deadlocks with Extended Events.
- ✅ Implement retry logic at the application level.

⚡ In short:

Deadlocks can't be completely eliminated in high-concurrency systems, but with **query design, indexing, and isolation level tuning**, you can reduce their occurrence significantly.

● SQL Server Deadlock – Real-time Scenarios (Q&A)

Scenario 1: Update Deadlock

Q: Two transactions update two tables but in different orders. Why does a deadlock occur, and how do you fix it?

A:

- **Cause:**
 - Transaction 1: Updates **TableA** → **TableB**
 - Transaction 2: Updates **TableB** → **TableA**
 - Both transactions wait on each other → deadlock.
- **Fix:**
 - Access resources in the **same order** across all transactions.
 - Example: Always update TableA before TableB.

Scenario 2: Reader vs Writer Deadlock

Q: A SELECT query (with shared locks) conflicts with an UPDATE query (exclusive lock). How to avoid deadlock?

A:

- **Cause:** A reader holds a shared lock while waiting for another row, and the writer waits for exclusive access.
- **Fix:**
 - Enable **READ COMMITTED SNAPSHOT ISOLATION (RCSI)** to use row versioning:
 - ALTER DATABASE MyDB SET READ_COMMITTED_SNAPSHOT ON;
 - Or use NOLOCK hints (with caution).

Scenario 3: Long Transaction Deadlock

Q: A large transaction updating thousands of rows causes frequent deadlocks. How do you fix it?

A:

- **Fix:**
 - Break large updates into **smaller batches** (e.g., 5000 rows at a time).
 - Commit frequently to release locks.
 - Ensure proper indexes to minimize locking.

Scenario 4: Missing Index Deadlock

Q: Why can missing indexes increase deadlocks?

A:

- **Cause:** SQL Server scans entire tables → acquires more locks → higher contention.
- **Fix:**
 - Create appropriate **covering indexes**.
 - Example: For frequent WHERE OrderDate > '2023-01-01', index on OrderDate.

Scenario 5: Deadlock in Stored Procedures

Q: Two stored procedures deadlock frequently. What's the best practice to avoid this?

A:

- **Fix:**
 - Ensure **procedures access objects in the same order**.
 - Apply consistent locking hints if necessary (UPDLOCK).
 - Keep transactions inside procedures **short** and **predictable**.

Scenario 6: Application-Level Prevention

Q: How should applications handle deadlock errors?

A:

- **Answer:**
 - Always catch SQL error **1205 (deadlock victim)**.
 - Implement **retry logic** (re-execute the transaction after a short delay).
 - Keep retries capped (e.g., 3 attempts).

Scenario 7: Reporting Query Causes Deadlock

Q: A heavy reporting query locks rows and blocks OLTP transactions, causing deadlocks. How do you fix it?

A:

- **Fix:**
 - Move reporting queries to a **read replica** or **read-only secondary** (AlwaysOn AG).
 - Or run reports under **SNAPSHOT** isolation.
 - Optimize reports with **indexed views** or **materialized aggregates**.

Scenario 8: Detecting Deadlocks

Q: How do you capture and analyze deadlocks in SQL Server?

A:

- **Options:**
 1. **Extended Events** → Capture xml_deadlock_report.
 2. **SQL Profiler** (older method).
 3. **Trace Flags:**
 4. DBCC TRACEON (1222, -1);
 5. Use **system_health session** (enabled by default).

Scenario 9: Deadlock Between Update & Select with Hints

Q: An UPDATE uses WITH (UPDLOCK), while a SELECT uses WITH (HOLDLOCK). They deadlock. Why?

A:

- **Cause:** Conflicting lock hints escalate contention.
- **Fix:**
 - Remove unnecessary lock hints.
 - Use proper isolation levels instead of forcing locks.

Scenario 10: Deadlocks in High-Concurrency System

Q: In a high-traffic OLTP system, deadlocks are frequent. What are global strategies?

A:

- Use **RCSI** to minimize read locks.
- Keep **transactions short**.
- Optimize schema with **proper indexing**.
- Ensure **consistent access patterns**.
- Add **retry logic** in the application.
- Monitor with **Extended Events** and fix hotspots.



Summary:

Deadlocks are not always avoidable, but they can be **minimized** by:

- Designing transactions and queries carefully.
- Using proper indexes and isolation levels.
- Monitoring deadlock graphs.
- Adding retry logic at the application layer.